A cache lookup costs under 0.4 % of the provider call

Median time per request, FAQ configurations; provider = Claude Haiku 4.5, 600 live calls

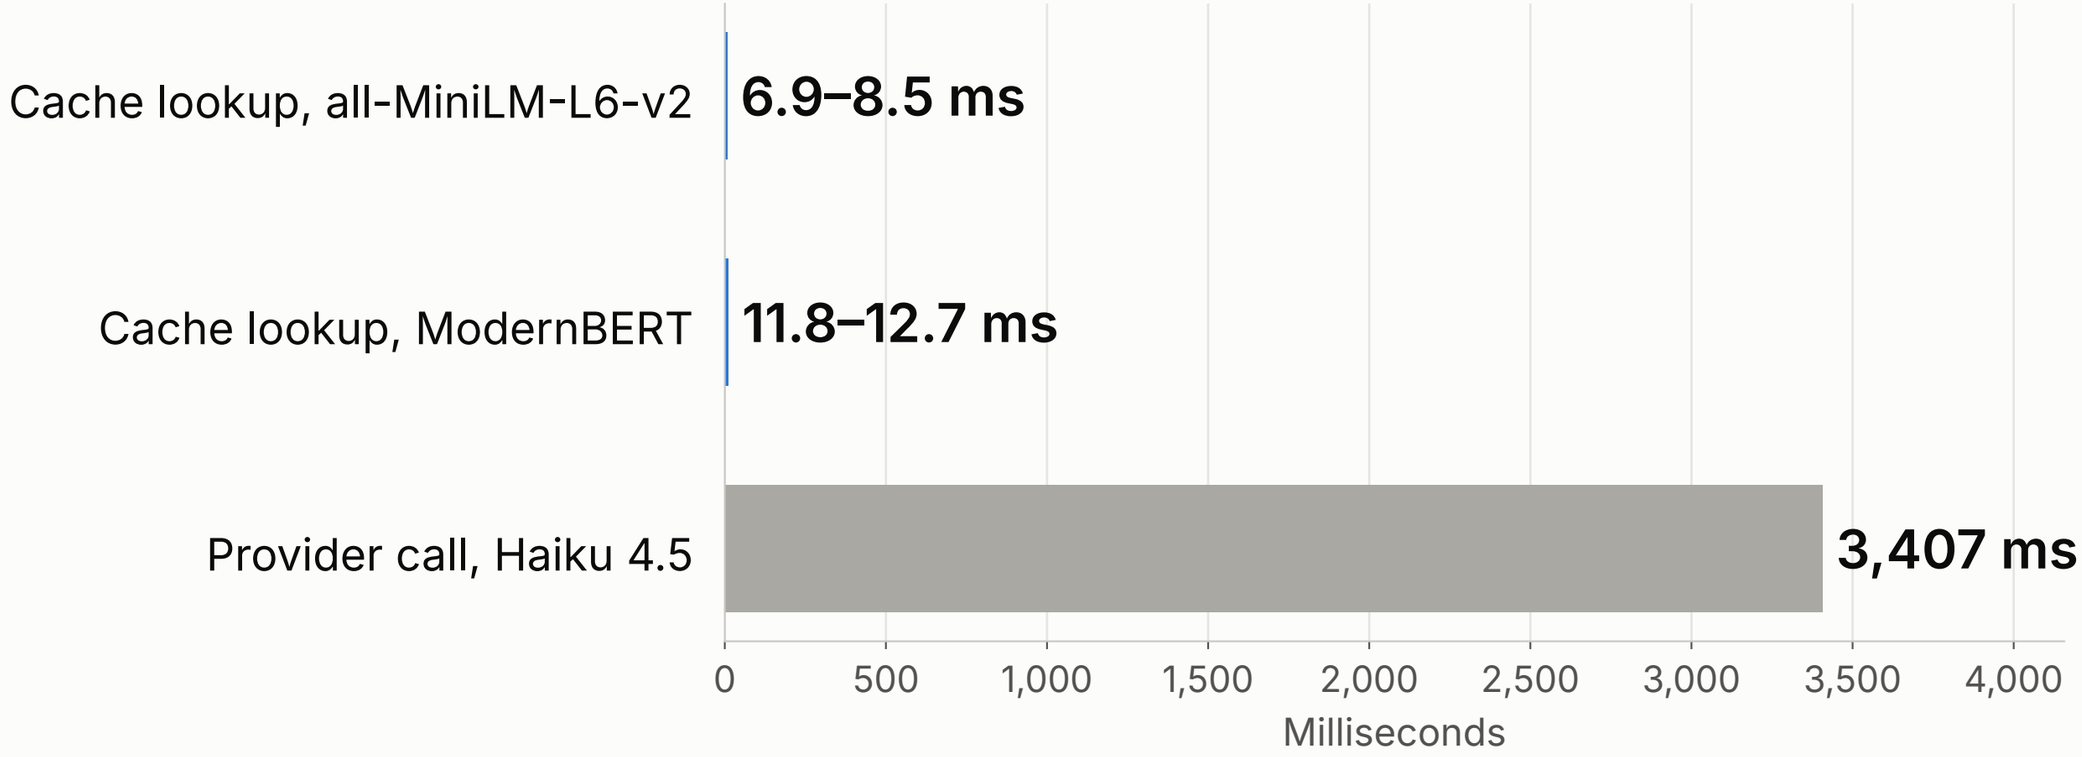

Break-even hit rate per request: 0.22 % (all-MiniLM-L6-v2) · 0.37 % (ModernBERT)

Every FAQ configuration clears it. Bars drawn to scale.

Source: release/latency/latency.csv (total_lookup_p50_ms), release/latency/llm_calls.json · dissertation §5.2.6