Cached requests run at up to 552 per second

Requests per second, all-hit vs all-miss · provider latency calibrated to Haiku 4.5 · no errors

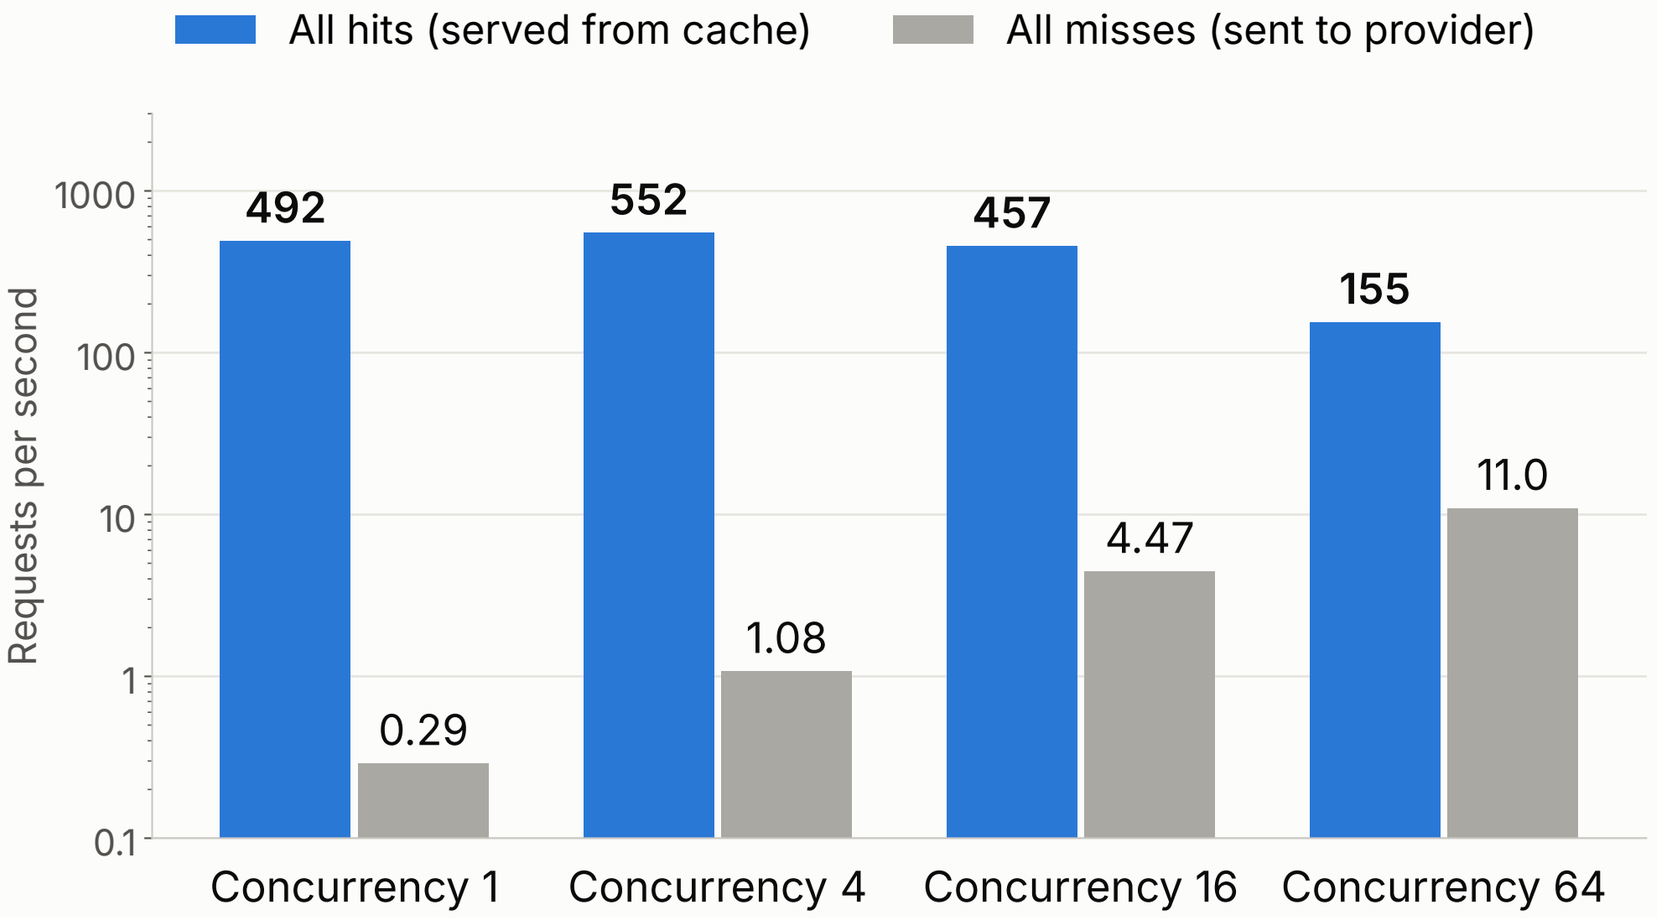

Hit vs miss throughput

Concurrency 1: 1,692×

Concurrency 4: 512×

Concurrency 16: 102×

Concurrency 64: 14×

Source: release/throughput/claude-haiku-4-5-20251001/throughput.csv · log scale · dissertation §5.2.6